

601.475 - Machine Learning

Kaushik Srinivasan

March 13, 2019

Topic 1: Introduction to ML	0	Topic 4: Gradient Descent	4
Topic 2: Linear Regression	1	Topic 5: Logistic Regression	5
Topic 3: Regularization	3		

Topic 1 — Introduction to ML

Machine Learning - Using *Experience* to gain *expertise*.
Design of algorithms of

- improve their performance
- at some task
- with experience

Supervised Learning

- **Classification** - Discrete Labels
- **Regression** - Continuous Labels

Task: Given $X \in \mathcal{X}$, predict $Y \in \mathcal{Y}$
Construct **prediction rule** $f : \mathcal{X} \rightarrow \mathcal{Y}$
Performance: Risk $R(f) \equiv \mathbb{E}_{XY}[\text{loss}(Y, f(X))]$,
 $(X, Y) \sim P_{XY}$
Experience: Training data $\{(X_i, Y_i)\}_{i=1}^n \sim P_{XY}$
(unknown)

$$\{(X_i, Y_i)\}_{i=1}^n \longrightarrow \text{Learning Algorithm} \longrightarrow \hat{f}_n$$

Unsupervised Learning

Task: Given $X \in \mathcal{X}$, learn $f(X)$

- Density Estimation
- Clustering
- Embedding

Performance Measure in Supervised L.

0/1 Loss: $\text{loss}(X, f(X)) = 1_{\{f(X) \neq Y\}}$

Square Loss : $\text{loss}(Y, f(X)) = (f(X) - Y)^2$

$\text{loss}(Y, f(X))$ - Measure of closeness between true label Y and prediction $f(X)$

$$(X, Y) \sim P_{XY}$$

Risk $R(f) \equiv \mathbb{E}_{XY}[\text{loss}(Y, f(X))]$

$\text{loss}(Y, f(X))$	Risk $R(f)$
$1_{f(X) \neq Y}$ 0/1 Loss	$P(f(X) \neq Y)$ Probability of Error
$(f(X) - Y)^2$ Square Loss	$\mathbb{E}[(f(X) - Y)^2]$ Mean Square Error

Bayes Optimal Rule

Ideal Goal: Construct **prediction rule** $f^* : \mathcal{X} \rightarrow \mathcal{Y}$

$$f^* = \arg \min_f \mathbb{E}_{XY}[\text{loss}(Y, f(X))]$$

Best possible performance:

$$\text{Bayes Risk} \quad R(f^*) \leq R(f) \text{ for all } f$$

Issues in ML

A good Machine Learning Algorithm:

- Does not *overfit* training data
- *Generalizes* well to test data.

Performance Revisited

Expected Risk (Generalization Error)

$$\mathbb{E}_{D_n} \left[R(\hat{f}_n) \right] \equiv \mathbb{E}_{D_n} \left[\mathbb{E}_{XY} [\text{loss}(Y, \hat{f}_n(X))] \right]$$

Ideal Goal: Construct **prediction rule** $f^* : \mathcal{X} \rightarrow \mathcal{Y}$

$$f^* = \arg \min_f \mathbb{E}_{XY}[\text{loss}(Y, f(X))]$$

Practical Goal: Given $\{X_i, Y_i\}_{i=1}^n$, $\hat{f}_n : \mathcal{X} \rightarrow \mathcal{Y}$, **learn** prediction rule

$$\hat{f}_n = \arg \min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n [\text{loss}(Y_i, f(X_i))]$$

By Law of Large Numbers: ($n \rightarrow \infty$)

$$\frac{1}{n} \sum_{i=1}^n [\text{loss}(Y_i, f(X_i))] \longrightarrow \mathbb{E}_{XY}[\text{loss}(Y, f(X))]$$

Consistency and Rates of Convergence

Excess Risk:

$$\mathbb{E}_{D_n} \left[R(\hat{f}_n) \right] - R(f^*)$$

is consistent if Excess Risk $\rightarrow 0$ as $n \rightarrow \infty$

How to Approach a ML Algorithm

1. Consider your Goal \rightarrow definition of task T .
2. Consider nature of experience E .
3. Choose type of output O to Learn
4. Choose performance measure P .
5. Choose representation for input X .
6. Choose set of possible solutions H .
7. Choose or design learning algorithm.

Topic 2 — Linear Regression

Formal setup

- Input data space \mathcal{X}
- Output (label, target) space \mathcal{Y}
- Unknown probability distribution $p(\cdot, \cdot)$ over $\mathcal{X} \times \mathcal{Y}$
- We are given labelled examples (\mathbf{x}_i, y_i) , $i = 1, \dots, N$ sampled i.i.d. from p ; $\mathbf{x}_i \in \mathcal{X}$, $y_i \in \mathcal{Y}$
- Goal: for any future \mathbf{x} , accurately predict y (drawn according to p) in other words: learn a mapping $f: \mathcal{X} \rightarrow \mathcal{Y}$

Types of Supervised Problems

Goal: learn $f: \mathcal{X} \rightarrow \mathcal{Y}$

- **Regression:** $\mathcal{Y} = \mathbb{R}$, learn (continuous) function f
- **Classification:** $\mathcal{Y} = \{1, \dots, C\}$, learn a separator between classes.

Linear Functions

General Form: $f(\mathbf{x}; \mathbf{w}) = w_0 + w_1 x_1 + \dots + w_d x_d$
Where $x_0 \equiv 1$.

- 1D case ($\mathcal{X} = \mathbb{R}$): a line
- $\mathcal{X} = \mathbb{R}^2$: a plane
- Hyperplane in general: d -D case

Loss Function

A **loss function**: $\ell: \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ maps prediction to cost, given true value. Standard choice of regression is squared loss: it's symmetric, non-negative, gives 0 loss for correct prediction.

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$

Empirical loss: LSQ minimizes empirical loss when ℓ is squared loss.

$$L(\mathbf{w}, \mathbf{X}, \mathbf{y}) = \frac{1}{N} \sum_{i=1}^N \ell(f(\mathbf{x}_i; \mathbf{w}), y_i)$$

Goal: Minimize the expected loss (**Risk**)

$$R(\mathbf{w}) = \mathbb{E}_{(\mathbf{x}_0, y_0) \sim p(\mathbf{x}, y)} [\ell(f(\mathbf{x}_0; \mathbf{w}), y_0)]$$

Empirical Risk Minimization (ERM) approach: to the extent that the training set is a representation of the underlying distribution $p(\mathbf{x}, y)$ the empirical loss serves as a proxy for the risk (expected loss). To minimize square loss -

$$f(\mathbf{x}; \mathbf{w}) = \mathbf{w} \cdot \mathbf{x}$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

How to we find $\mathbf{w}^* = [w_0^*, w_1^*, \dots, w_d^*]$

Least Squares

We need to minimize L w.r.t. \mathbf{w}

$$L(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

$$\frac{\partial}{\partial w_j} L(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \frac{\partial}{\partial w_j} (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2 = 0$$

$$= -\frac{2}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) x_{ij} = 0$$

$$\boxed{\sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) = 0}$$

Necessary Conditions:

1. Errors have zero mean
2. Errors are uncorrelated with the data..

Least Squares in Matrix Form

$$\mathbf{X} = \begin{bmatrix} 1 & x_{11} & \cdots & x_{1d} \\ \vdots & & & \\ 1 & x_{N1} & \cdots & x_{Nd} \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix}, \quad \mathbf{w} = \begin{bmatrix} w_0 \\ \vdots \\ w_d \end{bmatrix}$$

Prediction: $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$, errors: $\mathbf{y} - \mathbf{X}\mathbf{w}$

$$L(\mathbf{w}, \mathbf{X}, \mathbf{y}) = \frac{1}{N} (\mathbf{y} - \mathbf{X}\mathbf{w}) \cdot (\mathbf{y} - \mathbf{X}\mathbf{w}) = \frac{1}{N} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$

$$\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = \frac{1}{N} \frac{\partial}{\partial \mathbf{w}} [\mathbf{y}^T \mathbf{y} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{x}^T \mathbf{X}^T \mathbf{X} \mathbf{w}] = 0$$

$$= \frac{1}{N} [\mathbf{0} - (\mathbf{y}^T \mathbf{X})^T + 2\mathbf{X}^T \mathbf{X} \mathbf{w}] = 0$$

$$= -\frac{2}{N} (\mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{X} \mathbf{w}) = 0$$

$$\Rightarrow \boxed{\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}}$$

ML Paradox: The more training data we have, the "worse" the fit, but our prediction ability improves.

Best Unrestricted Predictor

$$f^* = \arg \min_{f: \mathcal{X} \rightarrow \mathbb{R}} \mathbb{E}_{(\mathbf{x}_0, y_0) \sim p(\mathbf{x}, y)} [(f(\mathbf{x}_0) - y)^2]$$

Chain rule of Probability: $p(\mathbf{x}, y) = p(y|\mathbf{x})p(\mathbf{x})$

By Def: $\mathbb{E}_{p(\mathbf{y}, \mathbf{x})} [g(y, \mathbf{x})] = \int_{\mathbf{x}} \int_y g(y, \mathbf{x}) p(y|\mathbf{x}) p(\mathbf{x}) dy d\mathbf{x}$

$$\mathbb{E}_{(\mathbf{x}_0, y_0) \sim p(\mathbf{x}, y)} [(f(\mathbf{x}_0) - y_0)^2]$$

$$= \mathbb{E}_{\mathbf{x}_0 \sim p(\mathbf{x})} [\mathbb{E}_{y_0 \sim p(y|\mathbf{x})} [(f(\mathbf{x}_0) - y)^2 | \mathbf{x}_0]]$$

$$= \int_{\mathbf{x}_0} \{ \mathbb{E}_{y_0 \sim p(y|\mathbf{x})} [(f(\mathbf{x}_0) - y_0 | \mathbf{x}_0)^2] \} p(\mathbf{x}_0) d\mathbf{x}_0$$

minimizing the inner conditional expectation for each \mathbf{x}_0

$$\frac{\partial}{\partial f(\mathbf{x})} \mathbb{E}_{p(y|\mathbf{x})} [(f(\mathbf{x}_0) - y_0 | \mathbf{x}_0)^2] = 2 \mathbb{E}_{p(y|\mathbf{x})} [f(\mathbf{x}_0) - y_0 | \mathbf{x}_0]$$

$$0 = 2(f(\mathbf{x}_0) - \mathbb{E}_{p(y|\mathbf{x})}[y_0 | \mathbf{x}_0]) \log p(Y|X; \mathbf{w}, \sigma) = \sum_{i=1}^N \log p(y_i | \mathbf{x}_i; \mathbf{w}, \sigma)$$

$$\hat{y}(\mathbf{x}_0) = f^*(\mathbf{x}_0) = \mathbb{E}_{p(y|\mathbf{x})}[y_0 | \mathbf{x}_0]$$

Generative vs. Discriminative Approach

Generative	Discriminative
<ul style="list-style-type: none"> Estimate Joint Probability Density $p(\mathbf{x}, y)$ <u>Normalize</u> to find conditional $p(y \mathbf{x})$ 	<ul style="list-style-type: none"> Estimate the density $p(y \mathbf{x})$ from data no need $p(\mathbf{x}, y)$

Decomposition of Error

We can understand the expected loss in this manner. $\hat{\mathbf{w}}$ are LSQ estimates from training data. \mathbf{w}^* are *optimal* linear regression parameters.


$$\begin{aligned} y - \hat{\mathbf{w}} \cdot \mathbf{x} &= (y - \mathbf{w}^* \cdot \mathbf{x}) + (\mathbf{w}^* \cdot \mathbf{x} - \hat{\mathbf{w}} \cdot \mathbf{x}) \\ \mathbb{E}_{p(\mathbf{x}, y)} [(y - \hat{\mathbf{w}} \cdot \mathbf{x})^2] &= \mathbb{E}_{p(\mathbf{x}, y)} [(y - \mathbf{w}^* \cdot \mathbf{x})^2] \\ &\quad + \mathbb{E}_{p(\mathbf{x}, y)} [(y - \mathbf{w}^* \cdot \mathbf{x})(\mathbf{w}^* \cdot \mathbf{x} - \hat{\mathbf{w}} \cdot \mathbf{x})] \\ &\quad + \mathbb{E}_{p(\mathbf{x}, y)} [(\mathbf{w}^* \cdot \mathbf{x} - \hat{\mathbf{w}} \cdot \mathbf{x})^2] \\ &= \mathbb{E}_{p(\mathbf{x}, y)} [(y - \mathbf{w}^* \cdot \mathbf{x})^2] \\ &\quad + \mathbb{E}_{p(\mathbf{x}, y)} [(\mathbf{w}^* \cdot \mathbf{x} - \hat{\mathbf{w}} \cdot \mathbf{x})^2] \end{aligned}$$

Approximation Error

$$\mathbb{E}[(y - \mathbf{w}^* \cdot \mathbf{x})^2]$$

Estimation Error

$$\mathbb{E}_{p(\mathbf{x}, y)} [(\mathbf{w}^* \cdot \mathbf{x} - \hat{\mathbf{w}} \cdot \mathbf{x})^2]$$



Statistical view of Regression

$$y = f(\mathbf{x}; \mathbf{w}) + \nu, \quad \nu \sim \mathcal{N}(\nu; 0, \sigma^2)$$

Where *noise* ν accounts everything not captured by f . Given input data \mathbf{x} , the label y is a random variable.

$$\begin{aligned} p(y|\mathbf{x}; \mathbf{w}, \sigma) &= \mathcal{N}(y; f(\mathbf{x}; \mathbf{w}), \sigma^2) \\ &= \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y - f(\mathbf{x}; \mathbf{w}))^2}{2\sigma^2}\right) \end{aligned}$$

To sample y for a given \mathbf{x}

Maximum Likelihood Estimation

$$\begin{aligned} \hat{\mathbf{w}}_{ML} &= \arg \max_{\mathbf{w}} p(Y|X; \mathbf{w}, \sigma) \\ &= \arg \max_{\mathbf{w}} \prod_{i=1}^N \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y_i - f(\mathbf{x}_i; \mathbf{w}))^2}{2\sigma^2}\right) \end{aligned}$$

We take the log likelihood because it is also monotonically increasing, and easier to use.

$$\begin{aligned} \log p(Y|X; \mathbf{w}, \sigma) &= \sum_{i=1}^N \log p(y_i | \mathbf{x}_i; \mathbf{w}, \sigma) \\ &= \sum_{i=1}^N \left[-\frac{(y_i - f(\mathbf{x}_i; \mathbf{w}))^2}{2\sigma^2} - \log \sigma\sqrt{2\pi} \right] \\ &= -\frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 - N \log \sigma\sqrt{2\pi} \end{aligned}$$

Red terms are independent of \mathbf{w} . Now we define *log-loss* as the negative conditional density of training data.

$$L(f(\mathbf{x}; \mathbf{w}), y) = -\log p(y|\mathbf{x}; \mathbf{w}, \sigma)$$

Maximizing log likelihood is always equivalent to *minimizing* log loss.

$$\begin{aligned} \arg \max_{\mathbf{w}} \sum_{i=1}^N \log p(y_i | \mathbf{x}_i; \mathbf{w}, \sigma) &= \arg \max_{\mathbf{w}} - \sum_{i=1}^N (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 \\ &= \arg \min_{\mathbf{w}} \sum_{i=1}^N (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 \end{aligned}$$

Maximum a Posteriori Estimation

Now given a prior $p(\theta)$ about our parameters, the maximum posterior is

$$\hat{\theta}_{MAP} = \arg \max_{\theta} p(\theta|X) = \arg \max_{\theta} \frac{p(X|\theta)p(\theta)}{p(X)}$$

The choice of prior matters! Bayesian approach, Utilitarian approach & regularization. If uniform $p(\theta)$, then $MAP = MLE$

Polynomial Regression

Consider the 1-D case where $f: x \rightarrow y$

$$f(x; \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_mx^m$$

No longer linear in x but linear in \mathbf{w} ! We define $\phi(x) = [1, x, x^2, \dots, x^m]^T$, then $f(x; \mathbf{w}) = \mathbf{w} \cdot \phi(x)$. So least squares solution:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \text{ where } \mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^m \\ 1 & x_2 & x_2^2 & \dots & x_2^m \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_N & x_N^2 & \dots & x_N^m \end{bmatrix}$$

$$f(\mathbf{x}; \mathbf{w}) = w_0\phi_0(\mathbf{x}) + w_1\phi_1(\mathbf{x}) + \dots + w_m\phi_m(\mathbf{x}) = \mathbf{w} \cdot \phi(\mathbf{x})$$


where $\phi_j(\mathbf{x}) : \mathcal{X} \rightarrow \mathbb{R}$, $j = 1 \dots, m$ are basis functions.

$$f(\mathbf{x}; \mathbf{w}) = \mathbf{w} \cdot \phi(\mathbf{x})$$

Is still linear in \mathbf{w} even when ϕ is non-linear in inputs \mathbf{x} ,
 $\phi_0(\mathbf{x}) \equiv 0$

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$\mathbf{X} = \begin{bmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \phi_2(\mathbf{x}_1) & \dots & \phi_m(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \phi_2(\mathbf{x}_2) & \dots & \phi_m(\mathbf{x}_2) \\ \dots & \dots & \dots & \dots & \dots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \phi_2(\mathbf{x}_N) & \dots & \phi_m(\mathbf{x}_N) \end{bmatrix}$$



How to avoid overfitting

- if a model overfits (it is *too* sensitive to data) it will be unstable - removing part of the data will change the fit significantly.
- We can *hold out* part of the data - this is *validation* set
- Fit the model to the rest and test on heldout data.
- Problem - if heldout set too small, we are susceptible to chance
- Problem - if heldout set too large, we get pessimistic (training on too little data compared to what we do).

Topic 3 — Regularization

Controlling for overfitting

1. More complex model (10th degree) overfits more than simple model (linear)
2. Pure ERM would always prefer complex model
3. Validation is a way to control for this in *model selection*

Intuitively, complexity of model measured by the number of “degrees of freedom” → more complex model more likely to overfit. Caused by finite training data.

Controlling for overfitting

1. **Idea 1:** Restrict model complexity based on amount of data → ≈ 10 examples per parameter

2. **Idea 2:** Directly penalize by number of parameters. Akaike information criterion (AIC).

$$\max \left[\log p(X | \hat{\mathbf{w}}) - \# \text{params} \right]$$

But: definition of model complexity as number of parameters is too simplistic.

3. **Idea 3:** consider the behavior of values of \mathbf{w}^*

Intuition. Should penalize not the parameters, but the number of bits required to encode the parameters → with finite parameter values, these are the same.

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \left\{ \frac{1}{2} \sum_{i=1}^N \log p(\text{data}_i; \mathbf{w}) - \text{penalty}(\mathbf{w}) \right\}$$

Ridge Regression (L_2)

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \left\{ \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2 + \lambda \sum_{j=1}^m w_j^2 \right\}$$

where $\mathbf{w} = [w_0, w_1, w_2, \dots, w_m]$. **Careful!** Solution is *not invariant* to scaling, we should normalize input before solving.

$$\mathbf{w}_{\text{ridge}}^* = (\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

We can write the following as an optimization problem

$$\min_{\mathbf{w}} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

subject to $\sum_{j=1}^m w_j^2 \leq t$

Lasso Regression (L_1)

$$\mathbf{w}_{\text{lasso}}^* = \arg \max_{\mathbf{w}} \left\{ - \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2 - \lambda \sum_{j=1}^m |w_j| \right\}$$

- Still concave (has a unique maximum), but not smooth (differentiable)
- Can solve using convex programming methods
- “lasso” → least absolute shrinkage and selection operator.


We can write the following as an optimization problem

$$\min_{\mathbf{w}} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

subject to $\sum_{j=1}^m |w_j| \leq t$


Geometry of surfaces

We can compare the shape of penalty as a function of w_j




An equivalent formulation of L_p regularization is constrained maximization.

$$\hat{\mathbf{w}} = \arg \max_{\mathbf{w}: \sum_{j=1}^m |w_j|^p \leq \beta} - \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$



- sufficiently large λ (small β) lasso leads to *sparsity*
- must explicitly solve optimization problem using Lagrange multipliers

Choice of λ



Most often λ is chosen by cross-validation.

Topic 4 — Gradient Descent

Often times, we cannot always solve the closed form solution $\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$ e.g. matrix is too large/difficult to calculate the pseudoinverse. **Gradient ascent/descent**

- Gradient ascent - “hill climbing” on function surface.
- start at a random spot and make steps in direction of maximal altitude.
- Equivalent: *gradient descent* on *convex* loss $-\log p(y | \mathbf{x}; \mathbf{w})$

Gradient descent algorithm

- Iteration counter $t = 0$
- Initialize $\mathbf{w}^{(t)}$ (to zero or a small random vector)
- for $t = 1, \dots$: compute gradient

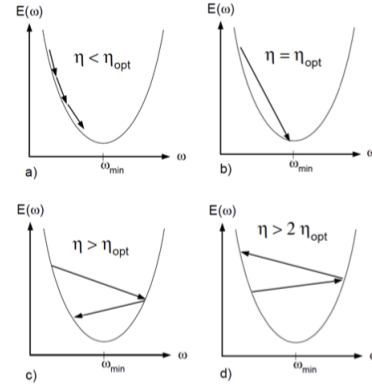
$$\mathbf{g}^{(t)} = \nabla f(\mathbf{X}, \mathbf{y}; \mathbf{w}^{(t-1)})$$

update model

$$\mathbf{w}^{(t)} = \mathbf{w}^{(t-1)} - \eta \mathbf{g}^{(t)}$$

- *learning rate* η controls step size

The size of η is important for the rate of convergence



Gradient descent convergence

- Minimum number of iterations (time budget) \rightarrow may not have converged
- Minimum required change in objective value (loss)
- Minimum required change in model parameters (\mathbf{w})

Estimation Theory

- *Estimator* $\hat{\theta}$ of parameter θ estimates it given input data.
- The *bias* of an estimator $\hat{\theta}$ is

$$\text{bias}(\hat{\theta}) \triangleq \mathbb{E}_X [\hat{\theta} - \theta]$$

- An *unbiased* estimator $\mathbb{E}_X [\hat{\theta}] = \theta$

Consistency of Estimator

- With enough data, bias *may* not be a problem.
- Estimator $\hat{\theta}$ is *consistent* if

$$\lim_{N \rightarrow \infty} \hat{\theta}_N \xrightarrow{p} \theta$$

Estimation and Regression

- True model $y = F(\mathbf{x}) + \nu$, 0 mean additive noise ν
- Approximate F by $f(\mathbf{x}; \hat{\mathbf{w}}) \in \mathcal{F}$ with $\hat{\mathbf{w}}$ estimated from data X .
- $\hat{f}(\mathbf{x}) = f(\mathbf{x}; \hat{w})$ estimate based on particular X
- $\bar{f}(\mathbf{x}) = \mathbb{E}_X [f(\mathbf{x}; \hat{\mathbf{w}})]$ average estimate over training sets X .

- $f^*(\mathbf{x}) = f(\mathbf{x}; \arg \min_w \mathbb{E}_{p(\mathbf{x}, y)}[(y - f(\mathbf{x}; \mathbf{w}))^2])$ the best estimate by function in \mathcal{F}

Bias-Variance decomposition

Denote $\bar{\theta} = \mathbb{E}[\hat{\theta}]$

$$\begin{aligned} E[(\hat{\theta} - \theta)^2] &= \mathbb{E}[(\hat{\theta} - \bar{\theta} + \bar{\theta} - \theta)^2] \\ &= \mathbb{E}[(\hat{\theta} - \bar{\theta})^2] + 2(\bar{\theta} - \theta) \underbrace{\mathbb{E}[\hat{\theta} - \bar{\theta}]}_{=0} + \mathbb{E}[(\bar{\theta} - \theta)^2] \\ &= \mathbb{E}[(\hat{\theta} - \bar{\theta})^2] - (\bar{\theta} - \theta)^2 \\ &= \text{var}(\hat{\theta}) + \text{bias}^2(\hat{\theta}) \end{aligned}$$

- **bias²** term \Leftrightarrow approximation error
- **variance** \Leftrightarrow estimation error due to finite data.

Bias-Variance in Regression

For a single \mathbf{x}_0

$$\mathbb{E}_X[(y_0 - \hat{f}(\mathbf{x}_0))^2] = (y_0 - \bar{f}(\mathbf{x}_0))^2 + \underbrace{\mathbb{E}_X[(\hat{f}(\mathbf{x}_0) - \bar{f}(\mathbf{x}_0))^2]}_{\text{variance}}$$

The first term can be further decomposed

$$(y_0 - \bar{f}(\mathbf{x}_0))^2 = \underbrace{(y_0 - F(\mathbf{x}_0))^2}_{\text{noise}} + \underbrace{(F(\mathbf{x}_0) - \hat{f}(\mathbf{x}_0))^2}_{\text{bias}^2}$$

- **noise** term is *irreducible*
- **bias²** is due to difference between f and F

$$\mathbb{E}[\text{squared loss}] = \text{bias}^2 + \text{var} + \text{noise}$$

Bias-Variance Tradeoff: Theory

Cramer-Rao inequality: for an unbiased estimator $\hat{\theta}_N$

$$\text{var}(\hat{\theta}_N) \geq \frac{1}{\mathbb{E}[(\frac{\partial}{\partial \theta} \log p(\mathbf{X}; \theta))^2]} = \frac{1}{\mathcal{I}(\theta)}$$

where $\mathcal{I}(\theta)$ is the *Fischer Information*. Intuitively, it measures the amount of data X provides about parameter θ

Topic 5 — Logistic Regression

Classification as Regression


$$f(\mathbf{x}; \hat{\mathbf{w}}) = w_0 + \hat{\mathbf{w}} \cdot \mathbf{x}$$

Can't just take $\hat{y} = f(\mathbf{x}; \hat{\mathbf{w}})$ as it won't be valid label. Hence we use a decision rule for $y \in \{-1, 1\}$

$$y = \begin{cases} 1 & f(\mathbf{x}; \hat{\mathbf{w}}) \geq 0 \\ -1 & f(\mathbf{x}; \hat{\mathbf{w}}) < 0 \end{cases}$$

$\hat{y} = \text{sign}(w_0 + \hat{\mathbf{w}} \cdot \mathbf{x})$ is also a valid *linear classifier* which transforms $\mathbb{R} \rightarrow \{1, -1\}$. Linear equation $w_0 + \hat{\mathbf{w}} \cdot \mathbf{x} = 0$ separates the space into two “half-spaces”

Classification as Regression



- $\mathbf{w} \cdot \mathbf{x} = 0$: a line passing through the origin and *orthogonal* to \mathbf{w}
- $\mathbf{w} \cdot \mathbf{x} + w_0 = 0$ shifts the line along \mathbf{w} .
- \mathbf{x}' is the projection of \mathbf{x} on \mathbf{w}
- Set up new coordinate system $\mathbf{x} \rightarrow (w_0 + \mathbf{w} \cdot \mathbf{x}) / \|\mathbf{w}\|$

Linear Classifiers

$$\hat{y} = h(\mathbf{x}) = \text{sign}(w_0 + \mathbf{w} \cdot \mathbf{x})$$

- Classifying using linear decision boundary effectively reduces data dimension to 1.
- Want to minimize the expected 0/1 loss for classifier $h : \mathcal{X} \rightarrow \mathcal{Y}$, which for (\mathbf{x}, y) is

$$L(h(\mathbf{x}), y) = \begin{cases} 0 & \text{if } h(\mathbf{x}) = y \\ 1 & \text{if } h(\mathbf{x}) \neq y \end{cases}$$

Risk of a classifier

- The risk (expected loss) of a C -way classifier $h(\mathbf{x})$

$$\begin{aligned} R(h) &= \mathbb{E}_{\mathbf{x}, y}[L(h(\mathbf{x}), y)] \\ &= \int_{\mathbf{x}} \sum_{c=1}^C L(h(\mathbf{x}), y) p(\mathbf{x}, y = c) d\mathbf{x} \\ &= \int_{\mathbf{x}} \left[\sum_{c=1}^C L(h(\mathbf{x}), y) p(y = c | \mathbf{x}) \right] p(\mathbf{x}) d\mathbf{x} \end{aligned}$$

- It is enough to minimize the *conditional risk* for any

\mathbf{x}

$$\begin{aligned} R(h|\mathbf{x}) &= \sum_{c=1}^C L(h(\mathbf{x}), c) p(y = c|\mathbf{x}) \\ &= 0 \cdot p(y = h(\mathbf{x})|\mathbf{x}) + 1 \cdot \sum_{c \neq h(\mathbf{x})} p(y = c|\mathbf{x}) \\ &= \sum_{c \neq h(\mathbf{x})} p(y = c|\mathbf{x}) = 1 - p(y = h(\mathbf{x})|\mathbf{x}) \end{aligned}$$

- To minimize conditional risk given \mathbf{x} , the classifier must decide

$$h(\mathbf{x}) = \arg \max_c p(y = c|\mathbf{x})$$

- This is the *best possible classifier* in terms of generalization i.e. expected misclassification rate on new examples.

Logistic Model

Define the decision boundary directly

$$\log \frac{p(y = 1|\mathbf{x})}{p(y = 0|\mathbf{x})} = w_0 + \mathbf{w} \cdot \mathbf{x} = 0$$


We define the logistic function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- $\sigma(0) = 1/2$. Shift the crossing to z : $\sigma(x - z)$
- change the “slope”: $\sigma(ax)$

Logistic Function in \mathbb{R}^d

- direction of \mathbf{w} determines orientation
- w_0 determines location
- $\|\mathbf{w}\|$ determines slope



Logistic Function in \mathbb{R}^d

- *Regression*: observe values, measure residuals under the model
- *Logistic Regression*: observe values, measure their probability under the model

$$\begin{aligned} p(y_i|\mathbf{x}_i; \mathbf{w}) &= \begin{cases} \sigma(w_0 + \mathbf{w} \cdot \mathbf{x}_i) & \text{if } y_i = 1 \\ 1 - \sigma(w_0 + \mathbf{w} \cdot \mathbf{x}_i) & \text{if } y_i = 0 \end{cases} \\ &= \sigma(w_0 + \mathbf{w} \cdot \mathbf{x}_i)^{y_i} (1 - \sigma(w_0 + \mathbf{w} \cdot \mathbf{x}_i))^{1-y_i} \end{aligned}$$

- the log likelihood of \mathbf{w}

$$\log p(Y|X; \mathbf{w}) = \sum_{i=1}^N \log p(y_i|\mathbf{x}_i; \mathbf{w})$$

- set derivatives to 0

$$\begin{aligned} \frac{\partial}{\partial w_0} \log p(Y|X; \mathbf{w}) &= \sum_{i=1}^N (y_i - \sigma(w_0 + \mathbf{w} \cdot \mathbf{x}_i)) = 0 \\ \frac{\partial}{\partial w_j} \log p(Y|X; \mathbf{w}) &= \sum_{i=1}^N (y_i - \sigma(w_0 + \mathbf{w} \cdot \mathbf{x}_i)) x_{ij} = 0 \end{aligned}$$

- Treat $y_i - p(y_i|\mathbf{x}_i) = y_i - \sigma(w_0 + \mathbf{w} \cdot \mathbf{x}_i)$ as *prediction error* of the model on \mathbf{x}_i, y_i